

# Regressão

November 10, 2023

```
[ ]: #import de todas as bibliotecas necessárias
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
[ ]: # Leitura do DataFrame
wine = pd.read_csv('winequalityN.csv')
#visualização das primeiras linhas do DataFrame
wine.head()
```

```
[ ]:
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	\
0	white	7.0	0.27	0.36	20.7	
1	white	6.3	0.30	0.34	1.6	
2	white	8.1	0.28	0.40	6.9	
3	white	7.2	0.23	0.32	8.5	
4	white	7.2	0.23	0.32	8.5	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	\
0	0.045	45.0	170.0	1.0010	3.00	
1	0.049	14.0	132.0	0.9940	3.30	
2	0.050	30.0	97.0	0.9951	3.26	
3	0.058	47.0	186.0	0.9956	3.19	
4	0.058	47.0	186.0	0.9956	3.19	

	sulphates	alcohol	quality
0	0.45	8.8	6
1	0.49	9.5	6
2	0.44	10.1	6
3	0.40	9.9	6
4	0.40	9.9	6

```
[ ]: wine.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   type                   6497 non-null   object
1   fixed acidity          6487 non-null   float64
2   volatile acidity       6489 non-null   float64
3   citric acid            6494 non-null   float64
4   residual sugar         6495 non-null   float64
5   chlorides              6495 non-null   float64
6   free sulfur dioxide    6497 non-null   float64
7   total sulfur dioxide   6497 non-null   float64
8   density                6497 non-null   float64
9   pH                    6488 non-null   float64
10  sulphates              6493 non-null   float64
11  alcohol                6497 non-null   float64
12  quality                6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB

```

```

[ ]: # Verificando se há valores nulos no DataFrame, retornando a soma de valores
      ↪ nulos por coluna
wine.dropna(inplace=True)
wine.isnull().sum()

```

```

[ ]: type                0
fixed acidity           0
volatile acidity        0
citric acid             0
residual sugar          0
chlorides               0
free sulfur dioxide     0
total sulfur dioxide    0
density                 0
pH                      0
sulphates               0
alcohol                 0
quality                 0
dtype: int64

```


```

[ ]: # Excluir colunas não numéricas
numeric_wine = wine.select_dtypes(exclude=['object'])

# Calcular correlação para as colunas numéricas
correlation_matrix = numeric_wine.corr()

```

```
# Plotar a matriz de correlação
plt.figure(figsize=(10,6))
sns.heatmap(correlation_matrix, annot=True)
plt.ylim(0, 12)
plt.show()
```



```
[ ]: #excluindo colunas com strings
wine = wine.drop(['type'], axis=1)
#separação dos dados em treino e teste
x_values = wine.drop(['alcohol'],axis=1).values
y_values = wine['alcohol'].values

x_train, x_test, y_train, y_test = train_test_split(x_values, y_values,
↳test_size=0.2, random_state=42)

# normalização
pipe = Pipeline([('scaler', StandardScaler())])
x_train = pipe.fit_transform(x_train)
x_test = pipe.transform(x_test)

# regressão (treinamento)
```

```
my_model = LinearRegression()
my_model.fit(x_train, y_train)
print(my_model.score(x_test, y_test)) # r2 squared
```

0.822867726346168

```
[ ]: #regressão (teste)
test_pred = my_model.predict(x_test)
train_pred = my_model.predict(x_train)

#função para avaliar o modelo
def print_evaluate(real, predicted):
    mae = metrics.mean_absolute_error(real, predicted)
    mse = metrics.mean_squared_error(real, predicted)
    rmse = np.sqrt(metrics.mean_squared_error(real, predicted))
    r2_square = metrics.r2_score(real, predicted)

    print('MAE:', mae)
    print('MSE:', mse)
    print('RMSE:', rmse)
    print('R2 Square', r2_square)
    print('-----')

print("Dados de teste:")
print_evaluate(y_test, test_pred)
print("\nDados de treinamento:")
print_evaluate(y_train, train_pred)
```

Dados de teste:

MAE: 0.3989843612973598

MSE: 0.26681814760634337

RMSE: 0.5165444294601805

R2 Square 0.822867726346168

-----

Dados de treinamento:

MAE: 0.3902429824443842

MSE: 0.2907334124101731

RMSE: 0.5391970070486047


R2 Square 0.7923181301457404

-----

```
[ ]: # Plotando os resultados
plt.scatter(y_test, test_pred)
plt.xlabel('True Values')
plt.ylabel('Predictions')
plt.title('True vs Predicted Values')
```

```
plt.show()
```

```
print(my_model.intercept_, my_model.coef_, my_model.score(x_test, y_test))
```



```
10.469660863957882 [ 0.67997065  0.26082324  0.06907571  0.93566117 -0.00667731
-0.00666862
-0.23469732 -1.72110406  0.45068922  0.19007688  0.12494332] 0.822867726346168
```